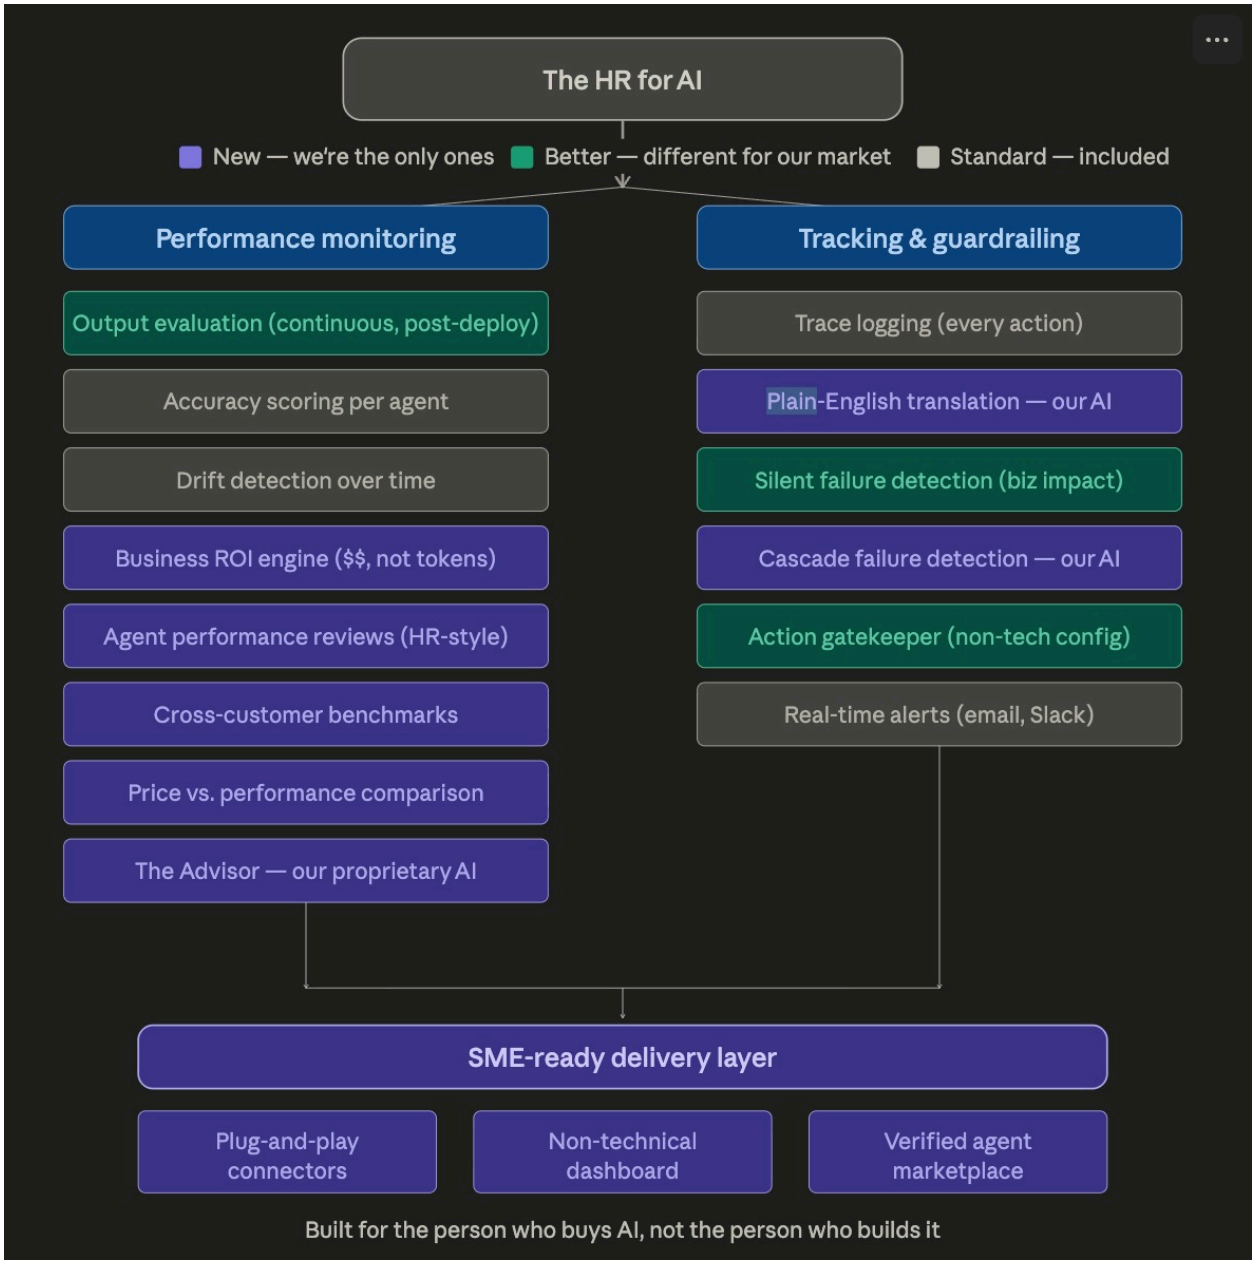
The HR for AI
New — we're the only ones
Better — different for our market
Standard — included
Performance monitoring
Output evaluation (continuous, post-deploy)
Accuracy scoring per agent
Drift detection over time
Business ROI engine ($$, not tokens)
Agent performance reviews (HR-style)
Cross-customer benchmarks
Price vs. performance comparison
The Advisor — our proprietary AI
Tracking & guardrailing
Trace logging (every action)
Plain-English translation — our AI
Silent failure detection (biz impact)
Cascade failure detection — our AI
Action gatekeeper (non-tech config)
Real-time alerts (email, Slack)
SME-ready delivery layer
Plug-and-play connectors
Non-technical dashboard
Verified agent marketplace
Built for the person who buys AI, not the person who builds it

ATLAS SYNAPSE — Product Feature Guide

The HR Department for Your AI

Performance Monitoring

How well are your AI agents doing their job, and are they worth the money?

Output Evaluation (Continuous, Post-Deploy)

🟢 BETTER — different for our market

What it does: Every time an agent responds, our AI runs a second-pass quality check on that response. Scores it for accuracy, relevance, and policy compliance. Flags bad outputs as incidents.

How it works:

- **Runs continuously** — not a one-time pre-deployment test. Checks live production traffic in real time.
- **LLM-as-judge** — sends the agent's input + output + context to our evaluation model. Returns pass/fail + confidence score + one-line reason.
- **Threshold-based flagging** — below 0.6 confidence = warning. Below 0.3 = critical. User can adjust later.
- **Operates on redacted data** — identity markers already stripped at the edge. Evaluator sees content, not customer identity.

Why it's different: Galileo and LangSmith do evaluation too, but they surface results as developer metrics in engineering consoles. We surface them as plain-language incident reports to the business owner.

When: MVP — Week 2-3. Proprietary AI, improves with every customer.

Accuracy Scoring Per Agent

⬜ STANDARD — included

What it does: Aggregates evaluation results into a single percentage per agent. "Your scheduling bot was 93% accurate this week." Headline number on the dashboard next to each agent.

How it works:

- **Simple math** — if 200 interactions and 186 passed evaluation, accuracy = 93%.
- **Failure count** — shows total failures in the last 7 days per agent.
- **Real-time** — updates as new traces are evaluated.

Why it's different: Standard analytics that competitors provide. We display it in business terms (% and count) instead of statistical distributions.

When: MVP — comes free from the evaluator data.

Drift Detection Over Time

⬜ STANDARD — included

What it does: Tracks each agent's accuracy, quality, and speed over weeks and months. Alerts the user when performance drops below their threshold. Catches silent degradation from model provider updates.

How it works:

- **Trend tracking** — plots accuracy and latency over time per agent.
- **Threshold alerts** — user sets a floor (e.g., 85% accuracy). If agent drops below, alert fires.
- **Catches model updates** — OpenAI/Anthropic push changes without notice. An agent at 94% last month might be 78% now. This catches it.

Why it's different: Arize and Galileo do drift detection for data scientists. We show it as "your agent is getting worse" with a plain explanation.

When: Phase 1.5 — needs 2-3 weeks of accumulated data before it's useful.

Business ROI Engine ($$, Not Tokens)

🟣 NEW — we're the only ones

What it does: Shows what each agent actually costs to run vs. what it saves. User inputs their baseline ("this task used to cost me X hours/month"). System calculates net value from trace data.

How it works:

- **Inputs** — user provides: what the task cost before AI (hours, dollars). One-time setup.
- **Outputs** — cost per decision, hours saved, error costs, agent running cost, net P&L per agent per month.
- **Business language** — "Agent A saved you $3,200 this month. Agent B cost more than it saved."

Why it's different: Every competitor tracks token costs. Nobody tracks whether the agent is worth employing. This is what makes a business owner renew.

When: MVP stretch (Week 6) or early Phase 2.

Agent Performance Reviews (HR-Style)

🟣 NEW — we're the only ones

What it does: Automated weekly or monthly report per agent, written in plain English by our AI. Covers: what it did, how well, what went wrong, what it cost, and whether it's improving or declining. Exportable PDF.

How it works:

- **Modeled on employee performance reviews** — a format every business owner already understands.
- **Generated by our AI** — uses accumulated trace data, evaluation scores, and cost data.
- **Exportable** — PDF download for sharing with partners, investors, or internal teams.

Why it's different: Completely unique. No competitor generates agent "performance reviews."

When: Phase 2 — needs several weeks of accumulated data to be meaningful.

Cross-Customer Benchmarks

🟣 NEW — we're the only ones

What it does: Tells a user how their agent compares to similar agents across our platform. "Your scheduling agent is 72nd percentile. Top performers average 96% accuracy — yours is at 88%."

How it works:

- **Aggregated + anonymized** — no customer sees another customer's data. Only aggregate stats.
- **Gets more accurate with scale** — every new customer makes the benchmarks better. This is the data moat.
- **Actionable** — if yours is below average, paired with Advisor recommendations to fix it.

Why it's different: No competitor can offer this without the same customer scale. Requires our data volume to exist.

When: Phase 2 — unlocks after 20-30 active accounts.

Price vs. Performance Comparison

🟣 NEW — we're the only ones

What it does: Compares what each agent costs to run (model pricing, token usage) against what it delivers (accuracy, task completion). Flags overspend: "You could save $X/month by switching to a cheaper model with the same performance."

How it works:

- **Model-aware** — knows the pricing of GPT-4, GPT-4o-mini, Claude Sonnet, etc.
- **Performance-aware** — compares accuracy and quality scores across models.
- **Produces swap recommendations** — "Switch from GPT-4 to GPT-4o-mini. Performance stays the same. Save $400/month."

Why it's different: No tool frames this as a business-level cost optimization. We turn monitoring data into money saved.

When: Phase 2 — needs enough trace data to compare model performance reliably.

The Advisor — Our Proprietary AI

🟣 NEW — we're the only ones

What it does: Recommendation engine trained on cross-customer data. Tells the user what to change: switch models, adjust prompts, replace underperformers. Replaces the AI consultant.

How it works:

- **Powered by benchmark data** — sees patterns across all customers. "Agents that include order history in their prompt perform 23% better."
- **Specific recommendations** — "Switch this agent to a cheaper model." "This agent needs more context." "This agent is underperforming — consider replacing."
- **Gets smarter over time** — trained on growing customer data. A competitor would need our data volume AND time investment.

Why it's different: Strongest moat. Proprietary AI that improves with every customer. Cannot be replicated without our scale.

When: Phase 2 — requires sufficient cross-customer data.

Tracking & Guardrailing

What is the agent actually doing, and how do we control it?

Trace Logging (Every Action)

⬜ STANDARD — included

What it does: Records every action an AI agent takes in an immutable, append-only log. Every prompt, response, tool call, tool result, decision point — all captured with timestamps, token counts, and cost.

How it works:

- **Immutable** — the trace cannot be edited or deleted by the agent, the user, or us.
- **Structured** — parent trace → child steps. Multi-step workflows captured as ordered chains.
- **Selective PII redaction at the edge** — identity markers (emails, phones, SSNs, cards) stripped by Cloudflare Worker before reaching our database. Content (what was said, advised, quoted) preserved.
- **Encrypted** — AES-256 at rest, TLS 1.3 in transit. 90-day retention. Deletion on request.

Why it's different: LangSmith, Langfuse, Braintrust all capture traces. They show JSON span trees for developers. We store the same data and translate it through the next feature.

When: MVP — Week 1-2. Foundation for everything else.

Plain-English Translation — Our AI

🟣 NEW — we're the only ones

What it does: Takes raw trace data (API calls, tool invocations, JSON) and converts it into a 2-3 sentence summary a non-technical person can read. "Your bot booked a customer for Tuesday at 3pm despite a calendar conflict because it misread the timezone."

How it works:

- **Runs on redacted data** — identity markers already stripped. Translation describes WHAT went wrong, not WHO it happened to.
- **Incorporates evaluation** — if the evaluator flagged the output, the translation explains why in business terms.
- **Primary display** — trace list shows the summary first. Raw (redacted) data available on expand.
- **Quality bar** — 8 out of 10 summaries understandable by a non-technical tester. Tolerable, not perfect for MVP.

Why it's different: Nobody translates traces into business language. This is the single biggest differentiator. It's what makes the entire product usable by someone who doesn't write code.

When: MVP — Week 3-4. Proprietary AI.

Silent Failure Detection (Business Impact)

🟢 BETTER — different for our market

What it does: Catches when an agent confidently gives a wrong answer that doesn't throw any error. The agent thinks it succeeded. The customer got bad info. Nobody inside the company knows until we flag it.

How it works:

- **Uses evaluator scores** — when output scores below the confidence threshold, it's flagged as an "incident" even though the agent reported no error.
- **Business-impact framing** — "Your bot told a customer the wrong return policy at 2:14pm" not "hallucination score: 0.34."
- **Paired with alerts** — flagged incidents trigger email notifications with the plain-English translation attached.

Why it's different: Galileo detects hallucinations as scores in a developer console. We describe the business impact in language the owner understands.

When: MVP — Week 4-5.

Cascade Failure Detection — Our AI

🟣 NEW — we're the only ones

What it does: In multi-step agent workflows, one wrong output early on silently corrupts everything downstream. At 85% accuracy per step, a 10-step workflow only succeeds 20% of the time. This catches Step 2 errors before they break Step 5.

How it works:

- **Monitors each step** — compares outputs at each step against what the next step expects.
- **Flags before propagation** — catches the error before the chain finishes executing.
- **Proprietary** — trained on accumulated trace data showing where multi-step workflows typically break.

Why it's different: Nobody does chain-level failure prediction for non-technical users.

When: Phase 2 — requires basic evaluator to be proven first + enough trace data.

Action Gatekeeper (Non-Tech Config)

🟢 BETTER — different for our market

What it does: Lets the business owner set rules for what agents can do without asking permission. Simple toggle-based UI, no code.

How it works:

- **Read-only actions** (looking up info, checking calendar) → auto-approved.
- **Reversible changes** (updating a record, rescheduling) → logged and spot-checked.
- **Irreversible actions** (sending emails, processing refunds) → requires human approval before the agent proceeds.

Why it's different: Galileo Protect and Zenity do guardrailing for engineers in developer consoles. Ours is configurable by a salon owner or office manager.

When: Phase 2 — retention/safety feature, not acquisition.

Real-Time Alerts (Email, Slack)

⬜ STANDARD — included

What it does: When the evaluator flags a failure, the system sends an immediate notification with enough context to act on it without opening the dashboard.

How it works:

- **Email contains** — which agent, when, plain-English translation of what went wrong, severity, link to dashboard.
- **Deduplication** — same failure 10x in an hour = one email with a count, not 10 emails.
- **Preferences** — on/off per agent. Immediate for critical. Daily digest for warnings.

- **Incident log** — all flagged failures stored in dashboard even if email is missed.

Why it's different: Standard alerting, but the content is in business language, not developer language.

When: MVP — Week 4-5.

SME-Ready Delivery Layer

Everything above, packaged so a non-technical business owner can use it without an engineer.

Plug-and-Play Connectors

🟣 NEW — we're the only ones

What it does: Zero-code integration. User clicks to connect from inside the platform they already use. No SDK, no API keys, no developer.

How it works:

- **Path A: Zapier/Make native app** — user clicks "Add Atlas Synapse" from their Zapier dashboard. One-click authorize. Traces flow automatically.
- **Path B: OAuth login** — Plaid-style popup. Log into AI provider, grant read-only access, done.
- **Path C: Builder installs for clients** — agent builder integrates once, all their customers get monitoring.
- **No credentials stored by us** — Zapier vault handles Path A. OAuth tokens for Path B. Project tokens for Path C.

Why it's different: Every monitoring competitor requires SDK installation by an engineer. We require zero code from the business owner.

When: MVP — Week 1-2.

Non-Technical Dashboard

🟣 NEW — we're the only ones

What it does: One screen. All agents. Plain English. Under 5 minutes from signup to first trace.

How it works:

- **Agent overview** — all agents with status, accuracy %, failure count, last active, cost.
- **Agent detail** — recent traces as plain-English summaries with eval scores. Expand for raw (redacted) data.
- **Incident feed** — all flagged failures, filterable by agent and severity. First thing the user checks.
- **Data transparency section** — shows what we store and what redacted data looks like. Builds trust without SOC 2.
- **Loads under 3 seconds** — paginated traces. Works on mobile.

Why it's different: Every competitor dashboard shows developer metrics. Ours shows business outcomes in language anyone can read.

When: MVP — Week 1-5 iterative.

Verified Agent Marketplace

🟣 NEW — we're the only ones

What it does: Browse, compare, and deploy AI agents that have been tested through our evaluation pipeline. Every listed agent has a performance score, failure rate, and cost profile before you buy.

How it works:

- **Builder submission** — agent builders submit their agents for testing.
- **Verification pipeline** — we run them through our Evaluator. Only passing agents get listed.
- **Monitoring pre-attached** — buyer gets automatic tracking from day one.
- **Network effects** — builders list because buyers are here. Buyers come because agents are pre-verified. Second revenue stream (10-20% commission).

Why it's different: GPT Store exists but has zero performance data. We sell agents with verified accuracy scores and monitoring built in.

When: Phase 2 — needs monitoring product to exist first.